

Capstone Project

MLflow Anomaly Detection

Akash Deep

June 2026

Dataset: data/dataset/no and data/dataset/yes

Experiment: anomaly-detection

1. Executive Summary

This report documents a complete anomaly detection workflow built for the MLflow capstone assignment. Four model configurations were trained on a local binary image dataset, compared using accuracy, recall (both classes), and macro F1 score, and tracked in a single MLflow experiment. The best model Random Forest was registered, promoted through the model registry (@challenger @champion), and validated with production inference on the held-out test set.

Key finding: Random Forest achieved the highest anomaly recall (74.2%) without SMOTE. Applying SMOTE to XGBoost actually reduced performance on this small, high-dimensional dataset.

2. Introduction



Detecting anomalies in operational data is a common real-world ML challenge. Class imbalance makes evaluation difficult accuracy alone can be misleading when one class dominates. This project implements the required experimentation pipeline: stratified splitting, SMOTE on training data only, four comparable models, MLflow tracking, and model registry workflow.

Dataset: 193 images in two folders no/ (class 0, normal) and yes/ (class 1, anomaly). Images were resized to 32-32 grayscale and flattened into 1,024 numerical features for Logistic Regression, Random Forest, and XGBoost.

3. Dataset and Preprocessing

3.1 Class Distribution

Total: 193 images | Class 0 (no): 89 | Class 1 (yes): 104 | Ratio: ~54:46. The dataset is mildly imbalanced. SMOTE was applied on the training set as required.




3.2 Train-Test Split (70/30, Stratified)

- Training set: 135 samples
- Test set: 58 samples untouched until final evaluation

3.3 SMOTE on Training Set Only

Before SMOTE: Class 0 = 62, Class 1 = 73. After SMOTE: both classes = 73 (146 total training samples).



4. Experiment Results

4.1 Logistic Regression (Baseline)

Parameters: C=1, solver=liblinear, used_smote=False

Accuracy: 0.4483 | Recall C0: 0.3333 | Recall C1: 0.5484 | F1 Macro: 0.4376

Weakest performer. Linear model struggled with 1,024 pixel features on limited data.

4.2 Random Forest

Parameters: n_estimators=30, max_depth=3, used_smote=False

Accuracy: 0.6207 | Recall C0: 0.4815 | Recall C1: 0.7419 | F1 Macro: 0.6091

Best overall model. Highest recall on anomaly class critical for detection tasks.

4.3 XGBoost (Original Data)

Parameters: eval_metric=logloss, used_smote=False

Accuracy: 0.5690 | Recall C0: 0.5185 | Recall C1: 0.6129 | F1 Macro: 0.5657

Second best. Better than Logistic Regression but below Random Forest.

4.4 XGBoost with SMOTE

Parameters: eval_metric=logloss, used_smote=True


Accuracy: 0.5000 | Recall C0: 0.4815 | Recall C1: 0.5161 | F1 Macro: 0.4987

SMOTE hurt performance. Oversampling added noise on this small high-dimensional set.

5. Metric Comparison

Model	Accuracy	Recall C0	Recall C1	F1 Macro	SMOTE
Random Forest	0.6207	0.4815	0.7419	0.6091	No
XGBClassifier	0.5690	0.5185	0.6129	0.5657	No
Logistic Regression	0.4483	0.3333	0.5484	0.4376	No
XGBoost + SMOTE	0.5000	0.4815	0.5161	0.4987	Yes





6. Observations on Class Imbalance Handling

- Dataset imbalance is mild (~54:46), so sample size and feature dimension mattered more than extreme skew.
- SMOTE balanced the training set (6273 for class 0) but did not improve XGBoost test performance.
- Random Forest without SMOTE achieved the best anomaly recall tree ensembles handled pixel patterns better.
- Selection criterion: recall_class_1 first, then f1_score_macro appropriate when missing anomalies is costly.
 - For production image tasks, a CNN would likely outperform raw pixels; this workflow meets capstone MLflow requirements.


7. MLflow Tracking and Model Registry

All runs logged to experiment 'anomaly-detection' on <http://127.0.0.1:5000>. Each run recorded hyperparameters, four metrics, and model artifacts.

Registry Workflow

- Registered best model as anomaly-detector-xgb-smote v1
- Assigned @challenger alias
- Copied to anomaly-detection-prod v1
- Assigned @champion alias
- Loaded models:/anomaly-detection-prod@champion test metrics matched Random Forest

7.1 Experiments List View



7.2 Runs Comparison View

Run ID:	54e6fc8373e94737a5164d71f9e6fb7a	beace459cd8e4e848b735ff4c096f286	f41fbcf586a944bba6eb0bd47ec9eab0	95deb719070c48ba88f634db26a2ce17
Run Name:	XGBClassifier With SMOTE	XGBClassifier	Random Forest	Logistic Regression
Start Time:	06/21/2026, 08:55:13 PM	06/21/2026, 08:54:52 PM	06/21/2026, 08:54:41 PM	06/21/2026, 08:54:18 PM
End Time:	06/21/2026, 08:55:32 PM	06/21/2026, 08:55:12 PM	06/21/2026, 08:54:52 PM	06/21/2026, 08:54:41 PM
Duration:	19.9s	20.6s	10.5s	22.9s

Parameters	
☒ Show diff only	
C	1
eval_metric	logloss
max_depth	3
n_estimators	30
solver	liblinear
used_smote	True

Metrics	
☒ Show diff only	
accuracy	0.5
f1_score_macro	0.499
recall_class_0	0.481
recall_class_1	0.516

7.3 Best Run Detail Random Forest

Random Forest	
Overview	Model metrics
Description	
No description	
Details	
Created at	06/21/2026, 08:54:41 PM
Created by	hp
Experiment ID	594281079765530894
Status	Finished
Run ID	f41fbcf586a944bba6eb0bd47ec9eab0
Duration	10.5s
Datasets used	—
Tags	Add tags
Source	C:\Users\hp\Downloads\Capstone Project\venv\lib\site-packages\ipykernel_launcher.py
Logged models	sklearn
Registered models	anomaly-detection-prod
Registered prompts	—

7.4 Model Registry @challenger

Registered Models

Share and manage machine learning models. [Learn more](#)

Filter registered models by name or tags

Name	Latest version	Aliased versions	Created by	Last modified	Tags
anomaly-detection-prod	Version 1	@champion - Version 1		06/21/2026, 08:55:34 ...	—
anomaly-detector-xgb-smote	Version 1	@challenger - Version 1		06/21/2026, 08:55:34 ...	—

New model registry UI

25 / page

7.5 Production Model @champion

anomaly-detection-prod

Created Time: 06/21/2026, 08:55:34 PM Last Modified: 06/21/2026, 08:55:34 PM

Description Edit

Tags

Versions Compare

Version	Registered at	Created by	Tags	Aliases	Description
Version 1	06/21/2026, 08:55:34 PM	@champion	Add		

1

8. Production Model Justification

Production model: Random Forest deployed as anomaly-detection-prod @champion.

Justification: Random Forest ranked first on both primary selection metrics recall on class 1 (0.7419) and macro F1 (0.6091). It correctly identifies ~74% of anomalies on the test set while maintaining the highest overall balance across classes. Production inference confirmed identical metrics, validating the end-to-end MLflow registry pipeline.

9. Conclusion

Four experiments were successfully tracked, compared, and registered in MLflow. Random Forest was selected for production. SMOTE was valuable to document but not beneficial for XGBoost on this dataset. The workflow demonstrates reproducible experimentation, metric-driven model selection, and production-ready model deployment through MLflow model registry.

10. References

- Dataset: Local repository data/dataset/no/, data/dataset/yes/
- MLflow: <https://mlflow.org/docs/latest/>
- imbalanced-learn SMOTE: <https://imbalanced-learn.org/>